

SEMANTIC MINIMALISM AND SEMANTIC INTERNALISM

EMMA BORG AND JOHN COLLINS

I — EMMA BORG

MUST A SEMANTIC MINIMALIST BE A SEMANTIC INTERNALIST?

I aim to show that a semantic minimalist need not also be a semantic internalist. §I introduces minimalism and internalism and argues that there is a *prima facie* case for a minimalist being an internalist. §II sketches some positive arguments for internalism which, if successful, show that a minimalist *must* be an internalist. §III goes on to reject these arguments and contends that the *prima facie* case for uniting minimalism and internalism is also not compelling. §IV returns to an objection from §I and argues for a way to meet it which does not depend on giving up semantic externalism.

The question addressed in this paper is whether a semantic minimalist must also be a semantic internalist. The key notions of minimalism and internalism will be explored below, but, in brief: according to minimal semantics semantic content attaches to sentences, it is determined exclusively by the syntactic form and lexical content of those sentences and it is maximally free from pragmatic effects. Also, in at least some versions of minimalism, a semantic theory is intended to capture the knowledge cognized by actual subjects which underpins their linguistic abilities. Now, all these claims find echoes within the approach of semantic internalism, which adds to them the claim that semantic content can be specified without appeal to the speaker's physical environment or linguistic community. The fact then that these two approaches share so many fundamental assumptions might lead us to expect minimalism to be sympathetic to internalism. Furthermore, it seems that the minimalist who endorses internalism can avoid an objection to her account (raised by contextualists such as Recanati 2004), to the effect that the content minimalism delivers is too minimal to count as genuine semantic content. So the conclusion of §I will be that there are *prima facie* reasons to hitch the minimalist wagon to the internalist horse. Then, in §II, I'll turn to look at some arguments which purport to show that a minimalist *must* be an internalist, since internalism is the only approach

which coheres with the minimalist's assumption that semantics be scientifically tractable. However, in §III, I'll reconsider these positive arguments for internalism and suggest that they do not show that a minimalist must be an internalist. Furthermore, embracing internalism does not, contrary to the suggestion of §I, offer a genuine response to the objection that minimal content is not semantic content. So I'll argue that even the *prima facie* case for uniting minimalism and internalism can be rejected. Finally, in §IV, I'll turn to an alternative response to the contextualist challenge of §I, a response which seeks to maintain the minimalist's claim to deliver genuine semantic content but without giving up on semantic externalism. I'll thus conclude that, if this alternative response is feasible, a semantic minimalist need not and should not be a semantic internalist after all.¹

I

Minimal Semantics and Semantic Internalism. According to minimal semantics we should draw a firm distinction between semantic content (which attaches to sentences relativized to contexts of utterance) and speech act content, allowing that while pragmatic effects are endemic in delivering the latter, their role in determining the former is extremely limited. In fact, according to the minimalist, the only role pragmatic processes play en route to delivery of semantic content is to provide values for overtly context-sensitive syntactic elements, such as indexicals, demonstratives and tense markers. For the minimalist, then, everything found at the semantic level is contributed by something in the syntax of the sentence and the apparatus taking one from the syntactic to the semantic can be modelled formally.² Furthermore, at least some minimalists (Borg 2004) want to claim that it is because semantic content is formally respectable that its recovery in a subject can be underpinned by a dedicated semantics module in the mind of the language user (where the notion of a module is un-

¹ I should point out that I'm not going to do very much to motivate externalism in this paper—if you don't feel the pull of externalism before we start it's unlikely that anything I say will reel you in. The aim of the paper is rather to explore some arguments which seem to show that one theory I'm rather partial to, namely semantic minimalism, is incompatible with another theory I'm rather partial to, namely semantic externalism. The paper is thus heavily biased from the outset.

² Exactly what 'formal modelling' here amounts to is open to question; see Borg (2007) for discussion.

derstood in computational terms, following Fodor 1983).

This picture, of semantic analysis as underpinned by an internal cognitive module, as driven by syntax and as (largely) uncontaminated by pragmatic processes, echoes the kind of claims which Chomsky and others have made for semantic content. However, according to Chomsky, semantic content should be specified in entirely internal terms: a theory of meaning need look no further than the limits of the individual mind.³ According to the internalist, then, the content of the term ‘water’ could (at least in principle) be shared by myself, my twin in an XYZ world and by a brain in a vat. Contrary to Putman (1975), meanings are in the head—if meanings are concepts then they are concepts whose content is independent of the world (Chomsky 2000, pp. 160–1). On this approach the proper subject of semantic inquiry is the state of the agent’s internal language faculty, so the study of semantics (along with the study of syntax or grammar) becomes a branch of individual psychology. This is not to say that our language never makes contact with the world, but the referential properties of language emerge at the point of *use* not at the point of *meaning*.⁴ So while a speaker might use a word with a given linguistic content to pick out some particular portion of the external world this use is not essential to the word’s meaning what it does, the meaning was not caused by the object in the world it is used to talk about and the content will remain constant over different kinds of referential use.

So one question we might now ask is: given the shared outlook of the two approaches in certain other respects, should a minimalist also embrace Chomsky’s semantic internalism? One reason to think that minimalism should opt for internalism is that such a pairing appears to avoid a serious objection to minimalism from contextualism. According to contextualists, pragmatic processes are capable of acting *twice* en route to the complete interpretation of a speech act: they can act once in the familiar post-semantic way, to yield implicatures, but they can also act prior to the determination of semantic content and are in this guise relevant to fixing the literal meaning of an uttered sentence. In fact, according to the contextualist, pragmat-

³ As we will see below, it is not always transparent exactly what one is committed to in being an internalist (or an externalist for that matter); indeed part of the thrust of this paper will be that we need not understand these terms in the way that Chomsky et al. seem to.

⁴ As Chomsky (2000, p. 36) notes, ‘terms themselves do not refer, at least not if the term *refer* is used in its natural language sense; but people can use them to refer to things’.

ic processes are endemic within the semantic realm. To see why the contextualist thinks this is the case, consider the following examples:

- (1) Jill is ready.
- (2) The apple is red.
- (3) That is red.

Sentence (1) looks simply incomplete without the addition of information from the context of utterance, even though no such additional information is required by anything (obvious) in the syntax of the sentence. We only recover something truth-evaluable, it seems, once we know from the context of utterance what Jill is ready to do. Although the sense of incompleteness is less obvious in (2) still it seems that without further contextual information we have something which falls short of being truth-evaluable. Unless we know the way in which the apple is supposed to be red (to degree n on its skin, etc.), it seems impossible to judge for a given world whether it is such as to make the sentence true or not.⁵ Finally, with respect to (3), I've argued elsewhere (Borg 2004; MS) that a minimalist account of the semantic content of a sentence like (3) cannot incorporate substantial identifying (i.e. non-linguistic) information about the referent. The claim is instead that someone is capable of grasping the semantic content of an utterance of this sentence just in case they can think about the referent under the token-reflexive linguistic description 'the actual object referred to by the speaker with this token of "that"'.

However, if this is what is involved in grasping the semantics of (3) then it cannot be the case that semantic content alone enables us to tell, for some arbitrary world, whether it satisfies the content or not, since an agent may well not be able to non-linguistically identify the referent of the demonstrative expression. Yet, Recanati objects to such minimal truth conditions:

This move strikes me as an unacceptable weakening of the notion of truth condition. The central idea of truth-conditional semantics (as opposed to mere 'translational semantics') is the idea that, via truth, we connect words and the world. If we know the truth conditions of a sentence, we know *which state of affairs must hold for the sentence to*

⁵ Space prevents an exploration here of how a minimalist might respond to the challenges of incompleteness and inappropriateness posed by sentences like (1) and (2); for details, see Borg (2004, ch. 4), and Cappelen and Lepore (2005).

be true; and that means that *we are able to specify that state of affairs*. T-sentences display knowledge of truth conditions in that sense only if the right-hand side of the biconditional is *used*, that is, only if the necessary and sufficient condition which it states is transparent to the utterer of the T-sentence. If I say ‘*Oscar cuts the sun* is true iff Oscar cuts the sun’, without knowing what it is to ‘cut the sun’, then the T-sentence I utter no more counts as displaying knowledge of truth conditions than if I utter it without knowing who Oscar is (for example, if I use the name ‘Oscar’ deferentially, in such a way that the right-hand side is not really *used*, but involves some kind of mention). (Recanati, (2004, pp. 92–3))

The challenge to minimalism here seems to be that whatever we can recover prior to rich appeal to contextual information must fall short of genuine semantic content, for semantic content should guarantee that a subject who grasps that content *knows how a world satisfying that content will be* (in some substantial, not merely disquotational sense). Yet minimal contents do not come with this guarantee. So, for instance, we might expect that knowledge of genuine truth conditions would permit an agent to sort worlds effectively into those which satisfy and those which fail to satisfy a given condition, yet this (as we saw in the case of demonstratives above) is not something which minimal truth conditions ensure. However, there is perhaps a response to this objection available to the minimalist who endorses internalism, for it could be objected that Recanati’s challenge is premised on too externalist a conception of semantic content. To think that content constitutes *semantic* content only if grasp of it brings in its wake substantial, non-disquotational knowledge of worlds is, the response might go, to impose a condition on a semantic theory which it is constitutionally unable to meet, since a semantic theory is a theory of internally specifiable (i.e. world-independent, brain-in-a-vat-compatible) content and not a theory of word-world connections (with these emerging only at the point of language *use*). So if semantic minimalism is paired with semantic internalism Recanati’s challenge seems to be avoided and this yields a *prima facie* reason for thinking a minimalist *should* also be an internalist. However in the next section I’d like to turn to some stronger arguments which seem to show that a minimalist *must* be an internalist.

II

The Arguments for Semantic Internalism. To be an internalist about semantic content is, it seems, to claim that linguistic meaning per se is not world-involving. This, of course, is diametrically opposed to externalism about semantic content—the thesis that the meanings of some words are at least in part determined by the environment to which subjects belong. Now externalism appears to have the status of the pre-theoretical position here, for it seems simply obvious that the meaning of a name like ‘Ian Botham’ is fixed by the person in the world it picks out, and that the meaning of a predicate like ‘square’ is given by the property of squareness had by external objects. However, a moment’s reflection shows us that any kind of simple-minded account of words whereby they are meaningful thanks to their relationship to things in the world must be mistaken. For even if it seems natural to assume that some words (like names) acquire meaning from the things in the world they pick out, lots of the things we say simply don’t seem to be connected to aspects of the external world in this way at all. For instance, we don’t suppose that there must be unicorns or sakes in order for the expressions ‘unicorn’ and ‘Jill’s sake’ to be meaningful (Quine 1953; Hinzen 2007, p. 13).

What is more, the internalist argument continues, even the assumption that apparently uncontroversial referring terms like ‘London’ or ‘Barack Obama’ refer to objects in some mind-independent physical world begins to look problematic under pressure, for it is not obvious that the relata posited in the relevant referential axioms really exist (Stainton 2006 labels this ‘the radical argument from ontology’).⁶ Thus given a putative referential clause, such as:

(a) ‘London’ refers to London.

where the expression on the left-hand side is supposed to name a linguistic expression (in this case, an English word) and the expression on the right introduces a real-world entity (in this case, the city of London), there is, the internalist objects, no such word and no such object. Turning first to the linguistic side of the relation: why might we be sceptical that there are words like ‘London’? Well the first point to note is that the public languages to which these words

⁶ Discussion in the next two paragraphs follows Stainton (2006) very clear introduction of the issues.

are supposed to belong themselves look pretty suspect, for there are no clear individuation conditions for public languages. Instead the distinctions between languages and dialects are often vague and blurry. Why is it that we count the quite similar and often mutually intelligible things said in the different countries of Northern Europe as belonging to distinct languages (English, French, etc.) rather than treating them all as variants within a more general public language (Romance), yet we treat very different and mutually unintelligible systems like Mandarin and Cantonese as mere dialects of the more general language Chinese? It seems clear that what drives the individuation of a language is not mind-independent facts about objective states of affairs in the world but rather a complex mish-mash of socio-economic factors (the kinds of things which resist any purely scientific approach to understanding them).

The same sort of vagaries which beset language individuation also crop up at the level of words. Thus we can ask whether we should posit one word 'in' and allow that there are many different ways for something to be in something else, or whether we should posit many different words (one for each sense?) all of which happen to share an orthographic presentation? And we can ask whether there is one word pronounced 'fotoGRAFer' in India and 'foTAH-grafer' in Canada, or whether these constitute two distinct words with the same meaning?⁷ It seems plausible to think that answering these and similar questions will be a matter of assessing our aims and interests in categorizing one way or the other, rather than an attempt to limn some objective fact of the world. So it seems that we lack the kind of clearly individuated, public words which the left-hand side of the schema assumes.

Furthermore, the radical argument from ontology continues, we should be sceptical about the putative objects on the right-hand side of (a) as well, for there are no physical, mind-independent entities of the kind such clauses require. The problem can be highlighted by noting that a term like 'London' can be used to pick out many different facets of the city it is supposed to refer to, thus in some contexts 'London' picks out a physical location ('London is east of Oxford'), in some a governmental structure ('London has a mayor'), and in some its inhabitants ('London is growing'). Indeed we can even run together such different aspects apparently without contra-

⁷ The example is from Stainton (2006, pp. 918–19).

diction, as when we say that ‘London is an ugly city but it is well-run’. The problem is that there simply cannot be external, real-world objects which are capable of having all the properties the referent of ‘London’ is supposed to have. As Chomsky writes:

Such terms as *London* are used to talk about the actual world, but there neither are nor are believed to be things-in-the-world with the properties of the intricate modes of reference that a city name encapsulates. (Chomsky 2000, p. 37)⁸

This scepticism about the physical status of the ordinary objects of reference is reinforced by noting that object individuation is interest-relative. For instance, as Carnap (1937) observed, if we ask how many objects are in a given box it seems that there is no simple right or wrong answer to be given. Instead the answer we should give depends on the conceptual framework we are working with: if we are counting only ‘middle-sized dry goods’ then one number is appropriate, but if we use a different conceptual scheme, say one which posits existence for mereologies, then some quite other number may well be appropriate. Yet if we can only count objects by adopting a particular conceptual scheme then this once again undermines the idea that there are objective, mind-independent objects of reference simply sitting around in the world, quite independent of us, waiting to be called upon by our language.

Finally, we might note along with Sosa that to admit objects of reference as objective parts of the external world is to submit the world to a kind of ontological explosion. Thus Sosa notes that we might introduce the term ‘snowdiscall’ into our language to refer to a collection of snow which has a shape somewhere between a ball and a disc; so every snowball is a snowdiscall but not every snowdiscall is a snowball. Further, each snowball and snowdiscall in turn must be a distinct entity from the piece of snow which comprises it since they have different persistence conditions—squashing a piece of snow

⁸ An internalist avoids this aspect of the radical argument for ontology since the expression on the right-hand side of a clause like ‘[‘London’ refers to London]’ picks out not an external object but a world-independent concept. Thus as Hinzen (2007, p. 82) writes: ‘The conclusion here should be that London, while having uniquely physical and non-physical aspects, has *none* of them essentially: it remains stable and self-identical across changes in apparently *any* of the properties that we might predicate of it ... There simply is no external object that we could point to and claim: *this* object is London, and it is the referent of the word *London* no matter what predication it is a part of, and it determines the meaning and how we use it to refer ... The only thing that remains stable in perspectively different acts of reference to the ‘same thing’ is the concept we have of that thing, and that concept alone.’

will suffice for the destruction of a snowball but not the destruction of the piece of snow which constituted it. But now Sosa notes:

[T]here are infinitely many shapes S_1, S_2, \dots , between roundness and flatness of a piece of snow, and, for each i , having a shape between flatness and S_i would give the form of a distinctive kind of entity to be compared with snowballs and snowdiscalls. Whenever a piece of snow constitutes a snowball, therefore, it constitutes infinitely many entities all sharing its place with it. Under a broadly Aristotelian conception, therefore, the barest flutter of the smallest leaf hence creates and destroys infinitely many things, and ordinary reality suffers a sort of ‘explosion’. (Sosa 1993, p. 620)

So, the radical argument from ontology concludes, clauses like (a) cannot form the basis of a semantic theory as there are simply no objects available to play the role of the required relata.

While the radical argument from ontology claims that there are no such things as ordinary words or ordinary objects of reference for those words, there is also a second, more modest, argument for internalism available. According to this ‘moderate argument from ontology’ (Stainton 2006), while ‘London’ and London do exist, they exist only as mind-dependent entities. They are thus not the kind of objects externalism promised us as the ordinary objects of reference nor can they figure within a rigorous science of language. Thus we find Stainton writing:

[B]eing objective and ignoring interest-relative distinctions, the ‘scientific perspective’ cannot see entities whose individuation conditions inherently involve complex-human interests and purposes ... [G]ranting that what common sense ‘sees’ is perfectly real, we still arrive at the same conclusion ... that a comprehensive science of language cannot (and should not try to) describe relations of semantic reference, i.e. word–world relations. That is because the things which manage to *be*, on this more moderate view ... are nevertheless not real *in the right sort of way*. Hence they cannot be ‘seen’ from the scientific perspective. (Stainton 2009, pp. 925–6)

That this is Chomsky’s position seems evident when he writes:

It is not that ordinary discourse fails to talk about the world, or that the particulars it describes do not exist, or that the accounts are too imprecise. Rather, the categories used and principles involved need not

have even loose counterparts in naturalistic inquiry. (Chomsky 2000, pp. 138–9)

So, the moderate argument from ontology concludes that semantic content, if it is to be scientifically respectable, must be internally specified.

Finally, it seems that the explanatory burden facing a semantic theory can be specified in purely internal terms. For instance, an adequate semantic theory for English needs to explain (in Pietroski's 2005 example) why there is no reading of the sentence 'The Senator called an oilman from Texas' whereby the Senator (rather than the call or the oilman) is from Texas. Yet explaining this doesn't require a complicated theory about the way words and objects hook up but rather a theory about the internal structures and content which influence semantic interpretation. This leads Pietroski to suggest:

[M]eaning is less tightly connected to the truth (and ontology and alien interpretability) than a lot of work suggests; expressions have semantic properties; but these are intrinsic properties of expressions that constrain without determining the truth conditions of utterances. One can say that semantics is a species of syntax on this view. But that is not an objection. Given how form constrains meaning in natural languages, perhaps we should indeed replace the idea that semantic properties are *not* syntactic properties with a suitably expansive view of syntax. (Pietroski 2005, p. 296)

If the arguments of this section are correct then it seems not only that a semantic minimalist *should be* a semantic internalist (the conclusion of §1) but that she *must* be a semantic internalist. For either there are no such things as the words and objects the externalist perspective presumes, or words and objects do exist but are individuated via appeal to human interests and are thus not the right kind of things to figure in a science of language. On the other hand, however, it also seems that the move to internalist semantics might leave the explanatory burden of our semantic theory unchanged for much of the work our theory must do is already characterized in internal terms.⁹ So it seems that there are good reasons for the minimalist to embrace internalist semantics and with it the solution to the contex-

⁹ Further evidence of the internalist burden for semantics comes from work in lexical semantics on the syntactic distribution of expressions, e.g. Levin (1993). Exploring this work is beyond the scope of the current paper, however the issue is taken up again in Borg (MS).

tualist challenge that minimal content is not genuine semantic content. However, in the next section I'd like to reconsider the arguments presented thus far, for on closer inspection it seems that there remain externalist moves to be made.

III

Rejecting the Arguments for Semantic Internalism. Turning first to the internalist's objection that the referential picture can't apply across the board: it is not obvious that the externalist should feel particularly discomforted by this point. First, externalism, as stated above, is the theory that *for at least some expressions* what they mean is determined by features of the agent's environment, but this clearly leaves room for other types of expression in natural language whose meaning is determined in other ways.¹⁰ Second, as we will see below, we need to be clear about exactly what we take externalism to be claiming by the statement that words get their meaning from facts outside the agent. For at least some varieties of externalism might be applicable to all terms in a natural language. Specifically, if one were to opt for the kind of social externalism promoted by Burge (1979), then externalism might be thought to hold true for a majority of natural language terms, not merely those which intuitively appear 'world-directed'. So, without further support, worries about the limitations of externalism do little to damage the externalist cause. However further support is, of course, in the offing here in the form of the two internalist arguments from ontology, so let's turn to these now.

According to the radical argument from ontology there are no such things as public words or the ordinary objects of reference, so referential clauses stated in terms of such entities are destined to be empty. Yet we might ask what motivation the radical argument has for setting the standards for existence so high, for as Stainton notes:

[P]erfectly real objects can be quite hard to individuate/count, and can be norm-bound. They need not require a 'robust divide', but can rather be objectively different only in degree, with human interests settling the kind-divide between them. One could thus allow that there is such

¹⁰ Of course, an opponent might go on to argue that the way in which these non-externalist terms come to be meaningful could be extended to cover all terms in a natural language, in which case externalism would prove otiose, but showing this obviously requires much more work from the internalist.

a thing as English ... and that the nature of English and the words/sentences in it depend on a host of complex relations (political, military, historical, religious, etc.)—including even explicitly normative ones having to do with ‘correct speech’ ... [T]his does not make English and its elements unreal. (Stainton 2006, p. 921)

It seems that this kind of approach could also extend to the objects on the right-hand side of clauses like ‘[‘London’ refers to London]’, allowing that, in at least some cases, human interests and actions do have a role to play in individuating the ordinary objects of reference without this making them unreal. Thus it could be that, in part, what makes something a chair is its maker’s intention that it be used as a chair. Or again what makes something London is in part agreement amongst a community of speakers that a current object is the same as that previously called ‘London’. Yet it doesn’t follow from this that there are no such things as chairs or London. If this is right then the radical argument from ontology can simply be rejected on the grounds that it sets existence conditions for both words and objects too high.

Furthermore, in response to the Carnapian concern that object individuation is dependent on the kind of conceptual framework the counter adopts, it seems that an externalist might admit this point without it undermining their essentially realist world view.¹¹ For while *recognizing* objects may well be a perspective-relative matter this doesn’t mean that there is not an objective world underpinning that perspective. As Sosa notes:

[F]rom the fundamentally and ineliminably perspectival character of our thought it does not follow that reality itself is fundamentally perspectival. Everything that is true relative to a perspective and everything that is false relative to a perspective may be as it is as a necessary consequence of the absolute and unperspectival character of things [O]ur perspectival references and truths may be seen to derive necessarily from absolute and unperspectival reality. (Sosa 1993, p. 608)

Finally, in response to Sosa’s concern about ontological explosion, we need to be careful where we locate the point of detonation here.

¹¹ Following Davidson (1974), we might also reject the idea that Carnapian counters really have different conceptual schemes here, since everything statable in one scheme can be translated into the other and a statement and its translation must be true or false together (i.e. the same statement can’t be true relative to one conceptual scheme and false relative to another).

For though it is right that our *conception of reality* expands when we admit snowdiscalls and all the other possible objects of reference this is not the same as *reality itself* expanding.¹² Snowdiscalls exist because their grounds for existence do (i.e. collections of snow of the relevant shape) and when we come to recognize them what we see in the world expands but reality remains unchanged. It seems then that the radical argument from ontology, with its strong conclusion that both ‘London’ and London fail to exist, can be rejected by the externalist: public words and the ordinary objects of reference they pick out may be individuated with respect to human interests and beliefs but this doesn’t entail that they do not exist nor that there is no objective reality underlying our perspectival conceptions.

At this juncture, however, the internalist can turn to the moderate argument from ontology, pointing out that, while all this may be true, it provides no succour to the minimalist who wants to hold on to semantic externalism. For such perspectival, interest-relative objects are not what externalism promised us nor are they the kinds of things which can figure in a science of language. So, to the extent that minimalism holds semantic content to be scientifically tractable, externalism must be abandoned. Once again however it seems that an externalist can query this line of argument. For a start it’s simply not all that clear what externalism promised us at the outset concerning the ordinary objects of reference. It is true that in Putnam’s classic externalist thought experiment ‘water’ was held to refer to H_2O —a purely objective stuff picked out via the (non-interest relative) vocabulary of physics—but it is far from clear that this was supposed to be the model for an externalist explanation of all natural kind terms, let alone all expressions in natural language. For instance, on Burge’s social externalism, what matters for the meaning of the term ‘arthritis’ is the meaning assigned to this expression by experts in the community, viz. *a painful inflammation of the joints*. This is clearly an externalist proposal even though we have not specified the content of the expression in terms of the basic vocabulary of physics. Externalism holds for an expression *e* if *e*’s content is fixed (in part) via an appeal to facts about an agent’s environment, either their physical environment or the community of speakers to which they belong. Yet it seems that *this* could hold true whether the content thus fixed is spelt out via an objective language

¹² See also Allen (1998).

of science, like H_2O , or via talk of some more human-relative feature of reality like chairs, or even entirely abstract objects, like numbers. Externalism need not, it seems, be taken as simply co-extensive with some kind of brute physicalism. So objects individuated via appeal to human propensities could play a role in an externalist account of linguistic content, but could such an account still fall within the remit of *science*?

Chomsky is adamant that it couldn't:

To be an Intentional Realist, it would seem, is about as reasonable as being a Desk- or Sound-of-Language- or Cat- or Matter-Realist; not that there are no such things as desks, etc., but that in the domain where questions of realism arise in a serious way, in the context of the search for the laws of nature, objects are not conceived from the peculiar perspectives provided by the concepts of common sense. It is widely held that 'mentalist talk and mental entities should eventually lose their place in our attempts to describe and explain the world' (Burge 1992). True enough, but it is hard to see the significance of the doctrine, since the same holds true, uncontroversially, for 'physicalist talk and physical entities' (to whatever extent the 'mental'–'physical' distinction is intelligible). (Chomsky 2000, p. 21)

So must a genuine science of language be prohibited from talking about both beliefs and desires and chairs and tables? Well, a first point to note is that even if it is widely held that mentalistic talk must eventually disappear from a scientific explanation of the world, it is not universally held. Thus we might agree with Crane and Mellor (1990) that it remains an open question as to what kind of vocabulary a complete and final science might deploy, with assumptions that it will not deploy mentalistic terms or the categories of common sense perhaps being little more than prejudice at this stage. Secondly, it seems that an externalist might in fact accept Chomsky's claims about the requirements of science and the nature of the ordinary objects of reference *without* being forced to accept the conclusion that a science of language must be blind to common-sense categorizations. Specifically, it seems that we might agree with Chomsky that it is in part human interests which individuate (some of) the objects of common sense and allow that a purely scientific account of such human-relative individuation is impossible (since it would need to be a 'science of everything' as Chomsky says), and yet still hold that the vocabulary of common sense could play a role in a genuinely scientific explanation of linguistic abilities. The move

here would be to bracket the properties appealed to by commonsense categories separately from issues about what makes something instantiate this property (with this latter issue being a potentially non-scientific, interest-relative matter).

On this approach, a scientific study of semantic content would be required to deliver a genuinely explanatory and predictive theory which showed how complex surface behaviour (e.g. subjects' ability to acquire language given only limited evidence, to recover the literal meaning of an indefinite number of sentences, and to use language to communicate) was the result of a less complex underlying structure together with systematic rules for manipulating that structure. Thus an account which took the meanings of primitive expressions as basic and provided rules of composition for those terms (and a canonical method for delivering the meanings of complex expressions from those rules together with the meanings of primitives) would count as scientifically respectable on the current view, even if it incorporated axioms utilizing the categories of common sense. For the properties denoted by commonsense terms like 'desk' or 'chair' can be counted as perfectly respectable properties (i.e. one can be a realist about them) and this is all that is required for a scientific study of semantics. Of course there are murky questions lurking here about what makes something a desk or a chair and answering such questions may well require some kind of unsystematic, unscientific approach making reference to community norms and human interests, but these questions are not ones which get voiced or require answering from within the semantic domain. On this conception, then, a semantic theory is required to specify that 'London buses are red' means that *London buses are red*. It is not expected to tell us why 'red' means *red* and not *blue*, nor is it required to guarantee that every subject who grasps this semantic information knows what is required to make something a bus, nor must they be able to tell, for every object they encounter, whether or not it is a bus (a point returned to in §4).

The move here is analogous to one I have suggested elsewhere (Borg 2004; MS) for providing a minimal account of the semantic content of directly referential expressions, like demonstratives. For demonstratives I've argued that all that is required to grasp the semantic content of an utterance like 'That is red' is that a subject introduce a syntactically triggered singular concept relating to this token of 'that' which has as its content the object referred to by the

speaker (where which object this is is settled by appeal to the speaker's referential intentions). Furthermore, a subject should be deemed capable of doing this *even if* she can conceptualize of that content only under the token-reflexive description *the actual object referred to by the speaker with this token of 'that'*. Similarly, then, for understanding general terms, like 'chair', what is required is that the subject deploy a general concept of the relevant property (*chairhood*) and this is something she can do *without* engaging in questions about what makes something a chair. It follows, then, that there may be cases where semantic information alone is insufficient to allow a subject to decide, for some object, whether or not it instantiates the property in question (i.e. whether or not it is a chair) and answering this question may thus require a non-semantic (and probably non-scientific) inquiry.

It seems to me then that the argument that a scientific methodology demands we excise chairs and tables from our referential canon is not well-founded. The radical argument from ontology can be rejected as it sets the standards for existence claims too high, while the moderate argument from ontology can be rejected since it conflates questions of concept content with questions of how we identify objects as falling under those concepts. Instead it seems to me that we might allow reference and extension identification to remain as murky as we like without this preventing the commonsense properties we use for categorizing the world from entering into genuinely explanatory hypotheses about linguistic content. We might also think, contrary to the suggestions of the last section, that such externalist assumptions are required, i.e. that contrary to Pietroski the explanatory burden of semantics, while in part internalist, is also in part externalist. For instance, we might hold that the only way to make sense of certain kinds of semantic error is to assume an externalist perceptive. So when someone learns that 'contract' means *mutual agreement* rather than *written agreement*, they take themselves to be corrected about the meaning of the term, yet this behaviour only makes sense given the externalist perspective that what matters for word meaning can lie outside the individual.¹³

The case for a minimalist necessarily being an internalist is not

¹³ The example is from Higginbotham (2006). Note as well that if there is *any* external dimension to meaning then externalism must be right. For a properly nuanced externalist account could in principle explain all internally specified requirements on a semantic theory, while the converse does not hold.

then, I think, proven, but what about the claim of §I that a minimalist *should be* an internalist, since it provides her with an answer to one of the most serious challenges to minimalism, namely that her account fails to deliver genuinely *semantic* content? To recapitulate: the challenge of §I was that grasp of semantic content should yield substantial (non-disquotational) knowledge of which state of affairs would satisfy that content. So genuine semantic content should allow a subject to sort worlds into those which satisfy and those which fail to satisfy a given content. However it seemed no such sorting ability was entailed by grasp of minimal content. The proposed response for the minimalist, then, was to claim that semantic content was internal content: what minimalism provides is narrow content which still requires application to a world, via use, before truth-evaluable content and substantial world-knowledge is delivered. However the problem now is that it is quite unclear that a minimalist who takes this line in response to the contextualist challenge is really responding to, rather than capitulating to, the contextualist. For the main tenet of contextualism is that semantic (propositional or truth-conditional) content emerges only *following* pragmatic processing. That is to say, semantic content is a property of *utterances* not *sentences*—it is people that refer to things and thereby express complete contents, not the sentences that they utter. In response to this challenge the minimalist was advised to opt for internalism—to accept that minimal content doesn't yield rich, world-identifying, non-disquotational knowledge but to maintain that semantic content per se is not about identifying worlds at all (as opposed to the *use* of semantic content). Yet such a move seems to cede everything the contextualist contended. Contextualism is, on this picture, right to think that it is language in use which secures the world and to think that minimal content is something less than truth-conditional content. What internalism gives us then is not a minimalist answer to Recanati's challenge but a way to submerge minimalism beneath contextualism. Thus it seems to me that the minimalist who wants to preserve her identity in the face of Recanati's challenge is well advised not to take the internalist turn but to look for an alternative solution. Finally, then, I want briefly to sketch one such alternative solution.

IV

Minimal Content Can Be Truth-Conditional Content. I think a minimalist should ask the contextualist exactly what condition on genuine semantic content minimal content fails to meet. In this section I'll outline three possible candidates. First, perhaps a genuine truth condition, T, must guarantee that, for every possible world, w, someone who grasps T is capable of telling whether or not w satisfies T. However this condition seems far too strong, for it would seem that the only condition capable of meeting this sort of constraint would be a verification condition and we have no reason to expect that truth conditions must be verification conditions.¹⁴ Second, perhaps a genuine truth condition T must equip someone who grasps it with the ability to tell for some large number of possible worlds whether or not T is satisfied in those worlds. However there are two points here: first, this is a condition which minimal content seems to meet. For instance, take the sentence 'Some apple is red' as having the minimal content *Some apple is red*. Such a content will sort worlds into those where some apple is red in some way (so worlds which satisfy the truth condition) and worlds where every apple is not red in any respect (so which fail to satisfy the truth condition), and one might imagine that we have a pretty large number of possible worlds on each side here. Of course, it is also true to say that a contextually enriched content, such as *that apple is red to degree n on most of its skin* would yield a more fine-grained sorting of worlds—it would get less worlds satisfying and more worlds failing to satisfy the content—but it is hard to see why this difference in number should be taken to ground a distinction between the semantic and the non-semantic. Indeed, and this brings us to the second point, one might worry about the very idea of grounding a difference in kind, i.e. semantic versus non-semantic, on a difference of degree, i.e. the number of worlds it enables a subject to sort, since we would need strong independent reasons to draw a line anywhere on the degree axis and assume that it cleaves the semantic from the

¹⁴ Notice that even contextually enriched propositions seem likely to fail this condition; say, following Searle and Recanati, the semantic content of 'John cut the grass' is the contextually enriched proposition that *John cut the grass in the usual sense of cutting associated with grass*. Still we can ask: is that proposition true in a world where John left five blades uncut, or in a world where he used a lawnmower for most of the lawn but scissors for the edges? Such questions don't seem answered by our contextualist semantic content alone.

non-semantic. It seems to me then that neither of our first two proposals are feasible.

Perhaps instead, then, the proper condition on truth-conditional content is not one concerned with what it *tells* the subject (i.e. not one about the sort of world-sorting abilities it endows on someone who grasps it), perhaps instead the proper condition is concerned simply with whether or not it sorts worlds appropriately, regardless of whether or not a subject has the ability to track this sorting in every case. On this approach then a truth condition T must suffice to determine, for every possible world, whether T is true or not in that world. As it stands, this condition is too strong (for instance, there are issues with worlds where objects fail to exist, worlds where fundamental laws are very different to those of this world, and perhaps with vagueness). However the basic idea—that truth conditions must determine how the world must be in order to satisfy them—is one that a minimalist can embrace. Minimal truth conditions do sort worlds effectively even though a subject who grasps a minimal truth condition may have to do more work to *tell* which worlds are sorted in which way. So, for instance, if we take the minimal content of a sentence like ‘That is red’, as uttered on a given occasion, to be that *the actual object referred to by the speaker with this token of ‘that’ is red*, then it is obvious that, given semantic content alone, a subject will be unable to tell for some world whether it is one which satisfies the truth condition or not, since semantic content alone does not guarantee a non-linguistic way to identify that referent. However the semantic content alone does suffice to sort worlds into those which satisfy and those which fail to satisfy it.

To summarize: conditions based on what grasp of a truth condition ought to allow a subject to do (e.g. sort all worlds, sort most worlds, etc) will be ones which minimal content runs the risk of failing to meet. However the minimalist can afford, I think, to be pretty sanguine about this, since she antecedently insists on a minimal conception of what possession of a semantic theory alone will allow one to do. On the other hand, it seems that any plausible condition based on what a truth condition itself should do will be one which a minimal theory of content will be able to meet. Thus the objection from §1 that minimal content is not genuine semantic content can be rejected, for the conditions minimal content fails to meet are independently held by minimalism to be inappropriate constraints on a semantic theory *per se*.

V

Conclusion. There are arguments which seem to show that a semantic minimalist should also be a semantic internalist: first, the two approaches share a number of fundamental assumptions and a minimalist who is also an internalist appears to be able to obviate the contextualist challenge that minimal content is not semantic content. Second, there are reasons to think that internally specified content is the only form of content which meets the minimalist's demands for a scientifically respectable study of semantics. However I've tried to suggest that these arguments are mistaken. First, a minimalist who takes the internalist turn does not respond to, so much as capitulate to, contextualism. Second, the arguments concerning the requirements of science can be rejected, for though the categories of common sense may not figure in fundamental physics, this does not rule out the possibility of their figuring in a systematic, explanatory and computational study of semantics. Third, the contextualist objection can be disarmed since it seems that minimal content meets appropriate conditions on semantic content, even though it does not equip a subject with the kinds of abilities which contextualism takes to be granted by knowledge of semantics.¹⁵

<Please give full address>
 Philosophy Department
 University of Reading
 e.g.n.borg@rdg.ac.uk

REFERENCES

- Allen, Robert F. 1998: 'Absolutism vs. Relativism in Contemporary Ontology'. *Journal of Philosophical Research*, 23, pp. ???–???.
 Borg, Emma 2004: *Minimal Semantics*. Oxford: Clarendon Press.
 —MS: *Minimalism and Minds*.
 Burge, Tyler 1979: 'Individualism and the Mental'. *Midwest Studies in Philosophy*, 4, pp. 73–121.

¹⁵ I am grateful to audiences at the universities of Valladolid and Reading for feedback on this paper. Special thanks also to Jonathan Dancy and Ernie Lepore for detailed comments on a draft version. Work on this paper was made possible by an AHRC Research Leave Award and a Philip Leverhulme Prize.

- Cappelen, Herman, and Ernie Lepore 2005: *Insensitive Semantics: A Defense of Semantic Minimalism and Speech Act Pluralism*. Oxford: Blackwell.
- Carnap, Rudolf 1937: *The Logical Syntax of Language*. New York: Harcourt Brace.
- Chomsky, Noam 2000: *New Horizons in the Study of Language and Mind*. Cambridge: Cambridge University Press.
- Crane, Tim and D. H. Mellor 1990: 'There is No Question of Physicism'. *Mind*, 99, pp. 185–206.
- Davidson, Donald 1974: 'On the Very Idea of a Conceptual Scheme'. *Proceedings and Addresses of the American Philosophical Association*, 47. Reprinted in his *Inquiries into Truth and Interpretation*, pp. 183–98. Oxford: Oxford University Press, 1984.
- Fodor, Jerry 1983: *Modularity of Mind*. Cambridge, MA: MIT Press.
- Higginbotham, James 2006: 'Languages and Idiolects: Their Language and Ours'. In Ernest Lepore and Barry C. Smith (eds.), *The Oxford Handbook of Philosophy of Language*, pp. 140–50. Oxford: Oxford University Press.
- Hinzen, Wolfram 2007: *An Essay on Names and Truth*. Oxford: Oxford University Press.
- Levin, Beth 1993: *English Verb Classes and Alternations: A Preliminary Investigation*. Chicago: Chicago University Press.
- Pietroski, Paul M. 2005: 'Meaning before Truth'. In Gerhard Preyer and Georg Peter (eds.), *Contextualism in Philosophy: Knowledge, Meaning, and Truth*, pp. 255–302. Oxford: Oxford University Press.
- Putnam, Hilary 1975: 'The Meaning of "Meaning"'. In *Mind, Language and Reality: Philosophical Papers*, vol. 2: , pp. 215–71. Cambridge: Cambridge University Press.
- Quine, Willard van Orman 1953: *From a Logical Point of View*. Cambridge, MA: Harvard University Press.
- Recanati, François 2004: *Literal Meaning*. Cambridge: Cambridge University Press.
- Sosa, Ernest 1993: 'Putnam's Pragmatic Realism'. *Journal of Philosophy*, 90, pp. 605–26.
- Stainton, Robert J. 2006: 'Meaning and Reference: Some Chomskian Themes'. In Ernest Lepore and Barry C. Smith (eds.), *The Oxford Handbook of Philosophy of Language*, pp. 913–40. Oxford: Oxford University Press.

